Assignment 1: Extracting Parallelism

Due Sep 8, before the class

The purpose of this assignment is for you

- to develop insight about how to extract parallelism from simple codes.

- to recognize cases where some form of parallelism may not be correct

- to acknowledge that sometimes adding work is necessary

Note: Often when talking about algorithms, the weights of tasks are denoted with their complexity.
Note: 1 and 2 are exercise/warm up. 3, 4 and 5 are harder. All problems are independent.

1 Transform

Consider the transform function:

```
void transform (int* a, int* b, int n) {
  for (int i=0; i<n; ++i)
    b[i] = f(a[i]);
}
```

Question: Extract the dependencies. Assume the call to `f` cost O(1).
Question: What is the width? the critical path? the work?
Question: How does a schedule look like on P processors.

2 Reduce

Consider the reduce function:

```
template<typename T, typename op>
T reduce (T* array, size_t n) {
  T result = array[0];
  for (int i=1; i<n; ++i)
    result = op (result, array[i]);
  return result;
}
```

So if you define `T` as `int` and `op` as `sum`, it boils down to computing the sum of the array. You could use `op` as `max` and compute the maximum value of the array.

2.1   int, sum

Consider first the `int, sum` case which computes the sum of an array of integers.
Question: Extract the dependencies of this problem. What is the width? the critical path? the work?
Question: Noticing that the different loop iterations could execute in any order. Introduce a mutual exclusion clause on the dependency graph. Does that help?
Question: Assuming you have P processors, rewrite the code to introduce one local variable per processor to store partial computation. Extract the dependencies now. What is the width, critical path and work ?
Question: What does a schedule look like on P processors?

2.2   Variants

Question: Would these two parallel versions (with mutual exclusion and with local variable) be correct for `int, max`? Why?
Question: Would these two parallel versions (with mutual exclusion and with local variable) be correct for `string, concat`? Why?
Question: Would these two parallel versions (with mutual exclusion and with local variable) be correct for `float, sum`? Why?
Question: Would these two parallel versions (with mutual exclusion and with local variable) be correct for `float, max`? Why?

3   Find first

3.1   in an array

Question: Write a sequential algorithm that search a value `val` in an array `arr` of size `n` and return the position `pos` of the first location where `arr[pos] == val`. (and returns `n` otherwise.)
Question: What is the complexity of this algorithm? (as a function of `pos` and `n`).
Question: Can you make a parallel algorithm with $\theta(n)$ work? What is its critical path and width?
Question: Can you make a parallel algorithm with $\theta(pos)$ work? What is its critical path and width?

3.2   in a linked list

Question: What do you think about doing the same thing with a linked list?

4   Prefix sum

Prefixsum is the algorithm that computes $pr[i] = \sum_{j \leq i} arr[i]$ and often written sequentially:

```c
void prefixsum (int* arr, int n, int* pr) {
  pr[0] = arr[0];
  for (int i=1; i<n; ++i)
    pr[i] = pr[i-1] + arr[i];
}
```

Question: What is the structure of the dependency of prefixsum?
Question: How can you make it parallel? (Hint: you have to add work, a single pass on the array is not enough)

5   Merge Sort

Question: Recall the merge sort algorithm.
Question: Extract dependency on the merge sort algorithm. What is the critical path, work, and width? (Hint: instead of using loop iterations as a task, you can use function calls and function return as tasks. Think that merge sort is recursive.)
Question: How does the schedule of such an algorithm look like when P=4?
Question: Can you extract more parallelism? (Hint: increase the amount of work slightly)